

Exploring Data Transfer with ~~NoC~~Network-on-Chip based Many-Core Processors

Comment [Editor1]: Remark: Please check the journal guidelines for the preferred use of abbreviations in the title.

ABSTRACT

[TBD]

Keywords

Many-core; ~~Multi-core~~Multicore; Network-on-Chip

1. INTRODUCTION

The evolution of a next generation computing platform oriented toward multi/many cores is ~~becoming~~ unavoidable to satisfy ~~a~~the demand for increasing ~~demand-of~~ computation in conjunction with reasonable power consumption. ~~In recent years, computation amount~~ Recently, the computational ability of single core ~~processor-is-constantly-reaching~~processors has reached its limit, and ~~persistence~~thus the applicability of Moore's law [19] is unclear. Multi/many core architecture is ~~a major~~an important trend ~~of interest in the last~~recent years. ~~Integrating as it~~integrates cores, ~~they to~~ realize high-performance and general-purpose computing with low power consumption. ~~This~~Hence, high energy efficiency is ~~the~~a superior feature of multi/many core platforms.

Comment [Editor2]: Remark: Please check the journal guidelines for the preferred format of in-text citations. In general, they should appear in a ascending numerical order.

~~Above~~ As detailed in the preceding paragraph, the demand for increasing ~~demand-of~~ computation in conjunction with reasonable power consumption drives the need for multi/many core architecture in ~~many~~several domains. Real-time embedded systems are ~~one~~an example of ~~adapting in which~~ multi/many core platforms ~~because~~are adapted since they face increasing processing requirements. ~~In this domain, countless opportunities~~ Extant studies have examined numerous applications of multi/many cores platforms ~~are discussed~~ [4], [25], [22], [21], [5].

For example, automotive systems ~~contain~~involve various applications and ~~some of their are sometimes characterized by~~ demands for high-performance computing. Automotive applications are responsible for ~~many~~several control systems such as the powertrain, ~~the~~ chassis, ~~the~~ steering wheel, driver assistance, and user interface. Advanced driver assistance ~~system (ADAS) is becoming more and more complex~~ systems are characterized by increasing complexity and computational requirements and ~~is required~~necessitate intelligence such as an autonomous driving system. ~~Their complexity and computing requirements are continuously growing. Automotive systems require more computing resources. Meanwhile, it needs to realize extraordinary~~ Moreover, they also require significant energy efficiency and ~~reduce the costs~~cost reduction. Although modern automotive systems are ~~constructed~~composed of ~~a lot of~~several Electronic Control Unit (ECU) managing subsystems, there is a

Comment [Editor3]:
Tip: American-British Style
In American English, a comma (called serial or Oxford comma) is inserted before "and" in a series of three or more items.

~~trend~~shift from numerous scattered ECUs to hierarchical multi/many core Domain Controllers (DC). ~~Obviously, the~~ Evidently, a hybrid scenario is also ~~available. Combining~~ applicable. Specifically, DCs ~~combine~~ performance with low power consumption, ~~DCs to~~ realize high parallel processing and efficient data transport. ~~In addition~~ Furthermore, flexible scalability of a many-core platform ~~assists~~ facilitates development efficiency.

~~For another~~ A related example, ~~involves~~ avionics systems are that also ~~responsible for~~ include various applications and ~~a need of~~ require cost efficiency. ~~With highly integrated trend and increasing demand of processing, multi~~ Multi/many core platforms ~~have opportunities~~ can be applied in the avionics domain. ~~due to their highly integrated characteristics and increasing processing demands~~, Low power consumption of cores leads to lower cooling requirements, ~~which and this is the~~ a critical issue for avionics. Multi/many core platforms reduce ~~not only~~ energy consumption ~~but also as well as~~ the costs and weight. ~~The~~ Weight ~~reduction of weight~~ is the most important ~~result factor~~ factor for avionics. ~~Integration, and integration~~ with the a multi/many core platform reduces the amount of processing ~~boardboards~~ boards and cooling units. Thus, ~~utilizing~~ the utilization of an integrated platform ~~saves~~ results in space, weight, and ~~the costs~~ cost reductions.

~~Considering above background~~ Hence, multi/many core platforms are ~~drawn up~~ fabricated and released as commercial off-the-shelf (COTS) ~~multi-core~~ multicore components. ~~They are delivered significant attention in recent years. Kalray's the~~ Increasing research attention has focused on multi/many core platforms. The Multi-Purpose Processing Array (MPPA) ~~developed by Kalray 256 [12], [11], [15], Intel's~~ developed by Kalray 256 [12], [11], [15], Intel's Single-chip Cloud Computer (SCC) ~~developed by Intel [1],[3], and Tiler's~~ developed by Intel [1],[3], and Tiler's Tile64 ~~developed by Tiler [2] have clustered include the clustering of~~ developed by Tiler [2] have clustered include the clustering of many-core architectures, ~~where in which~~ cores are mapped closely. Their The clusters of cores are capable of performing separate independent applications with respect to the desired power envelope of embedded applications. Kalray's MPPA-256 packs 256 general-purpose cores, ~~which is overwhelming. This significantly exceeds the number of cores in other COTS's cores. With high performance per watt, COTS, and~~ it targets embedded systems, ~~HPC~~ high-performance computing (HPC), image processing, and networking, ~~due to the high performance per watt. Intel's~~ Xeon Phi [7] [8] is one of high performance computing (HPC) accelerators. an HPC accelerator. Xeon Phi is based on x86 architecture and ~~targets~~ use-case of data centers and workstation.

~~Although~~ Despite the emergence of the need ~~effor~~ for multi/many core platforms ~~has emerged, there are~~, several difficulties ~~for~~ persist in the adaptation of these platforms to real-time embedded systems [4], [25]. These difficulties are caused by ~~their~~ the hardware architecture and strict requirements of embedded systems. ~~One~~ A difficulty ~~is that cores share involves the sharing of~~ is that cores share involves the sharing of numerous resources (e.g., memory subsystems and I/O devices). ~~Since parallelized~~ by cores. Parallelized complex processes share some resources, ~~contention to these is frequently occurred, and this~~ leads to the frequent occurrence of conflicts. Cache coherency is also a critical problem ~~owing to their~~ caused by the presence of numerous cores. These difficulties disturb predictable timing behavior and software analysis. ~~Ensuring~~ Thus, the timing requirement of real-time embedded systems ~~is still an open issue. In multi/many core platforms, the~~ continues to warrant solutions. The impact of integrating real

Comment [Editor4]: Remark: Please consider revising this as "...its targeted use corresponds to data centers and workstations" if it retains your intended meaning.

time applications [in multi/many core platforms](#), is not completely understood ~~yet~~[to date](#). This is a critical issue because embedded systems have requirements for reliable and predictable behavior.

Contributions: ~~In this paper, we clarify the~~[This study focused on clarifying](#) performance characteristics of currently-achievable data transfer methods for many-core computing based on [Network-on-Chips \(NoC\)](#), such as MPPA-256, while ~~unveiling~~[examining](#) several new data transfer methods.

[TBD]

To the best of ~~our~~[the authors'](#) knowledge, this is the first ~~evidence of~~[study that examines](#) data transfer matters for many-core computing beyond an intuitive expectation, ~~which allows to allow~~ system designers to choose appropriate data transfer methods ~~depending based~~ on the requirement of their latency-sensitive many-core applications. ~~These~~[It is expected that the](#) findings [of the study](#) are ~~also~~[potentially](#) applicable for ~~some~~[several](#) types of many-core architectures ~~rather than as opposed to solely for~~ MPPA-256. ~~We believe that the~~[The](#) contributions of this ~~paper are useful for~~ [study can be used in](#) low-latency many-core computing.

Organization: The remainder of this ~~paper~~[study](#) is organized as follows. First, ~~Section 2 summarizes our considering the~~ system model ~~for considered in this paper. We present our~~ [study is discussed in Section 2 in which the](#) hardware model,

~~1. e. namely~~ Kalray MPPA-256 Bostan, and ~~our~~[the](#) system model ~~there are presented~~. Second, Section 4 illustrates ~~our~~ experimental evaluations. ~~Then~~[Subsequently](#), Section 5 ~~presents~~[examines](#) related ~~work about~~[studies that focus on](#) multi/many core systems. Finally, Section 6 ~~concludes this paper presents the conclusions and suggests directions for future work research~~.

2. SYSTEM MODEL

~~In this~~[This](#) section, ~~we present our~~ [presents the](#) system model used throughout ~~this paper. We consider the~~ [study. The](#) many-core model of Kalray MPPA-256 Bostan, ~~is considered~~. First, ~~the~~[a](#) hardware model is introduced in Section 2.1, ~~and this is~~ followed by ~~the~~[a](#) software model in Section 2.2.

2.1— Hardware Model

The MPPA-256 processor is based on an array of compute clusters (CCs) and I/O subsystems (IOSs) that are connected to nodes of Network-on-Chip (NoC) with a toroidal 2D topology (~~see as shown in~~ Figures 1, 2, and 3). The MPPA MANY-CORE chip integrates 16 ~~computing~~ clusters and 4 ~~I/O subsystems~~[IOS](#) on NoC. ~~We present the~~[The](#) architecture of Kalray MPPA-256 ~~is presented~~ in this section.

2.1.1 —~~I/O Subsystems~~[IOSs](#) (IOS)

Comment [Editor5]: Remark: Please confirm that this change is accurate.

MPPA-256 has the following four I/O subsystems (IOSes): North, South, East, and West IOSes. The North and South IOSes are connected to a DDR interface and an eight-lane PCIe controller. The East and West IOSes are connected to a quad 10Gb/s Ethernet controller. Two pairs of IOSes organize two I/O clusters (IOCs) as shown in Figure 1.

Each I/O subsystem (IOS) comprises of quad core resource managers (RMs) and a network interface.

Resource Managers (RMs): RM cores are connected to a 16-banked parallel shared memory of 2MB with a total capacity (IO SMEM) of 2 MB as shown in the left side of Figure 1. The four RMs have their own instruction cache 8-way associative of 32 (8 x 4) KB and share a data cache 8-way associative of 128 KB and external DDR access. Sharing the data cache of 128 KB makes it coherent, allowing coherency between the RMs. Additionally, RM cores operate controllers for the PCIe, Ethernet, Interlaken, and other I/O devices. They are able to operate the local peripherals, including the network interfaces with DMA. We can also make possible to conduct an application run on these RMs.




Figure 1: Overview of the architecture of the Kalray MPPA-256 Bostan architecture.

Network Interface: The network interface contains four DMA engines (DMA0-3) and four NoC routers as shown in the left side of Figure 1, and the IOS DMA engine manages

transfers between the IO SMEM, the IOS DDR, and

the IOS peripherals (e.g., PCIe interface and Ethernet controllers). Through NoC Routers, The DMA engine transfers data between routers on NoC, through NoC routers. The DMA engine has the following three NoC interfaces: a receive (Rx) interface, a transmit (Tx) interface, and a micro core (UC). A micro core is a network processor programmable that can be programmed to set threads sending data with a Tx interface. A UC is able to extract data from memory by using a programmed pattern and to send the data on the NoC. Once it is initiated, this is made continues in an autonomous fashion without using a Processing Elements (PE) and an RM.

2.1.2 Compute Clusters (CC)

In MPPA-256, the 16 inner nodes of the NoC correspond to the CCs. Figure 2 illustrates the architecture of each CC.

• **Processing Elements (PEs) and an RM:** In a CC, 16 processing elements (PEs) and ~~one~~^a RM share ~~2MB~~^{2 MB} cluster local memory (SMEM) ~~that is~~ composed of 16 independent memory banks ~~is~~^{with 16 independent memory banks of 16,384 x 64bit. Each bank of 64-bit.} ~~The capacity of each SMEM has a capacity of~~^{bank corresponds to} 128 KB. ~~These~~^{The} PEs are mainly used by users for parallel processing. Developers spawn computing threads on PEs. The PEs and an RM in CC ~~are~~^{correspond to} the Kalray-1 cores, which implement a 32-bit 5-issue Very Long Instruction Word (~~VLIW~~) architecture with

~~16W-600MHz~~^{16 W-600 MHz} (typical) or ~~24W-800MHz~~^{24 W-800 MHz}. Each core is

fitted with its own instruction and data caches. Each cache is ^a 2-way associative with a capacity of 8 KB. Thus, 17 k1-cores (a PE or the RM) share ^a multi-banked

~~2MB~~^{2 MB} SMEM.

• **~~A~~ Debug Support Unit (DSU) and ~~a~~ Network Interface:**

In addition to PEs and an RM, bus masters on SMEM ~~are the debug support unit~~ (~~correspond to a DSU~~) and a DMA engine in a network interface. A DMA engine and a NoC router are laid out in a network interface. ~~Similar~~^{In a manner similar} to IOS, the CC DMA engine ~~has~~ also ^{has the following} three interfaces: an Rx, a Tx, and a UC. It is ~~instantiated in every cluster and connected to the SMEM.~~

Comment [Editor6]: Remark: Please confirm that this change is accurate.

Comment [Editor7]: Remark: Note that the intended meaning is not clear. Please check this term and revise accordingly.

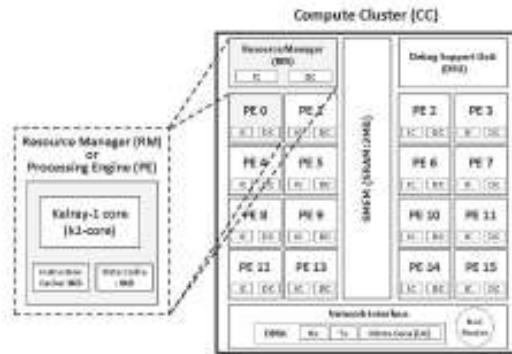


Figure 2: ~~Compute-Computation of Cluster (CC)~~ architecture.

2.1.3 ~~Network-on-Chip (NoC)~~

The 16 CCs and the 4 IOSs are connected by a NoC as shown in Figure 3. Furthermore, a NoC is constructed on the bus network and has routers on each node.

• **Bus Network:** Bus network connects nodes (CCs and IOSs) with a torus topology [9], which has ~~twice-dual~~ bands and ~~takes~~ involves a low average number of hops ~~when~~ compared to mesh topology [27], [26]. The network is actually composed of the following two parallel ~~NoC~~ NoCs with bi-directional links (denoted by red lines in Figure 3): the data NoC (D-NoC), ~~which~~ that is optimized for bulk data transfers; and the control NoC (C-NoC), ~~which~~ that is optimized for small messages at low latency. Functionally, NoC ~~is~~ corresponds to a packet switched network. Data are packaged in variable length packets ~~which~~ that circulate between routers in a wormhole manner: ~~in which~~ packets are broken into small pieces called flits (flow control digits). The NoC traffic is segmented into packets, ~~and~~ each packet ~~has~~ includes 1 to 4 header flits and 0 to 62 payload data flits.

Comment [Editor8]: Remark: Please confirm that this change is accurate.

• **NoC Routers:** A node per compute cluster and four nodes per I/O subsystem ~~hold~~ holds the following two ~~own~~ routers: a D-NoC router and a C-NoC router. Each RM on NoC node (CC or IOS) is associated with ~~these~~ the fore-mentioned two NoC routers. Furthermore, DMA engines in a network interface on the CC/IOS send and receive flits through ~~these~~ the D-NoC routers with the Rx interface, the Tx interface, and the UC. A mailbox component ~~which is~~ corresponds to the virtual interface for the C-NoC ~~and~~ enables one-to-one, N-to-one, or one-to-N low-latency synchronizations. The NoC routers shown in Figures 1 and 2 illustrate nodes as R0-15, R128-131,

Comment [Editor9]: Remark: Please consider revising this as "routers of its own."

R160-163, R224-227, and R192-195 in Figure 3). For

purposes of simplicity, we illustrate D-NoC/C-NoC routers are illustrated with one NoC router. In both D-NoC and C-NoC, each network node (CC or IOS) has includes the following 5-link NoC routers: four duplexed links for north/east/west and south neighbours, one neighbors and a duplexed link for local address space attached to the NoC Router-router. The NoC routers have include FIFOs queuing flits for each direction. The data links are four bytes wide in each direction and operate at the CPU clock rate of 600MHz/600 MHz or 800MHz/800 MHz, and therefore, each tile



Figure 3: Network-on-Chip (NoC) connections (both D-NoC

and C-NoC).

can transmit/receive a total of 2.4GB/4 GB/s or 3.2GB/2 GB/s, which is spread across the four directions (i.e., north, south, east, and west).

2.2 Software Model

A Figure 5 shows the software stack used for Kalray MPPA-245 used in this paper is shown in Figure 4. Kalray's MPPA the present study. The Kalray system is an extensible and scalable array of computing cores and memory. On such a With respect to this type of a system, we can it is possible to map several programming programming models or runtimes, e.g., such as Linux, real-time Operating System operating system, POSIX API, OpenCL, and OpenMP. We describe each Each layer is described in detail.

In ~~the~~ hardware abstraction layer, an abstraction package abstracts hardware of a CC, IOS, and NoC. ~~This~~ The abstraction package serves as a system that ~~doesn't~~ does not provide any services. The hardware abstraction is responsible for partitioning ~~the~~ hardware resources and controlling access to ~~these~~ the resources from the user-space operating system libraries. ~~Moreover~~ Additionally, the abstraction package ~~is able to~~ retrieve ~~the~~ retrieves resources allocated to a partition at any time. It ~~is able to set~~ sets-up and ~~control~~ controls inter-partition communications as a virtual machine abstraction layer. The hardware abstraction runs on the dedicated RM core. All the services are commonly provided by a rich operating system (~~for e.g.,~~ virtual memory, interrupts, ~~scheduler, etc.,~~ and schedule) that must be provided by user-space libraries. Consequently, each runtime or operating system implements its own ~~required~~-services that are optimized to ~~its~~ specific needs. This is because each ~~programming~~ programing model or runtime ~~has~~ involves different requirements. ~~Minimal~~ A minimal kernel avoids ~~the waste~~ wastage of resources and mismatched needs.

In a low-level ~~Library~~ library layer, ~~the~~ Kalray system also provides libraries for handling NoC. ~~Additionally,~~ NoC features, such as routing- and quality of service, ~~are to be~~ set by the programmer. The Libnoc allows direct access to memory mapped registers for their configurations and ~~use~~ uses. It is designed to cause a minimum amount of CPU overhead. It also serves as a minimal abstraction for



Figure 4: ~~Kalray MPPA-256's~~ The software stack of Kalray MPPA-256.

resource allocation. The Librouting offers ~~the~~ a minimal set of functions ~~that can be used to use-for-routing-route~~ data between any clusters of the MPPA, ~~both with including~~ unicast (one target) ~~mode~~ modes or multicast (multiple targets) ~~mode~~. Routing modes. As shown in Figure 3, routing on the torus network ~~shown in Figure 3~~ is statically conducted with its own policy. The Libpower enables spawning and waiting for the end of execution of a remote cluster.

~~In OS layer, various~~ Various operating systems support the abstraction package. ~~We introduce in the OS layer. The~~ following Real-Time

Operating System (RTOS) ~~is introduced~~:

- **RTEMS:** RTEMS, ~~the~~ (Real-Time Executive for Multiprocessor Systems) is a full featured RTOS prepared for embedded platforms. It supports several APIs and standards, and most notably supports the POSIX API. The system provides a rich set of features, and an RTEMS application ~~is most of the time just mostly corresponds to~~ a regular C or ~~C++~~ C program ~~using that uses~~ the POSIX API. ~~We are able to build~~ Additionally,

RTEMS can be built on the IOC.

- **NodeOS:** On CC, ~~we can build the MPPA cluster~~ MPPA cluster operating system utilizes a runtime called NodeOS. ~~This~~ NodeOS. The OS addresses the need for a multicore OS ~~conforming as much as to conform to the maximum possible extent~~ to the standard POSIX API. The NodeOS enables a user code by using POSIX API to run on PEs on CC. First, NodeOS runtime starts on PE0 ~~before~~ prior to calling the user main function. ~~Then, we can call~~ Subsequently, pthread is called on other PEs.

Comment [Editor10]: Remark:
Please consider revising this to "C++."

Comment [Editor11]: Remark:
Please check whether the edits retain your intended meaning.

• **eMCOS:** On both CC and IOS, eMCOS provides minimal ~~programming~~programming interfaces and libraries. ~~Specifically,~~ eMCOS is a real-time embedded operating system developed by eSOL, ~~(a Japanese supplier for RTOS), and~~ eMCOS is the world's first commercially available many-core RTOS for use in embedded systems. This OS implements a distributed micro-kernel architecture. ~~This~~The compact micro-kernel is equipped with only minimal functions. It enables applications to operate priority based message passing, local thread scheduling, and thread management on ~~not only~~IOS but also as well as CC.

RTMES and NodeOS are provided by Kalray, and eMCOS

is released by eSOL.

3. DATA TRANSFER FRAMEWORK

In this section, ~~we explain~~data transfer methods in MPPA-256 are explained. For scalability ~~purposes~~purposes, MPPA-256 accepts clustered architectures ~~wherein which~~each cluster contains its own memory. 16 cores are packed as a cluster and they share 2 MB memory (SMEM) as shown in Figure 2. ~~Avoiding~~This avoids frequent memory contention ~~by due to~~by due to numerous cores, ~~this and~~this and helps ~~increment of~~in increasing the number of cores. However, ~~this architectures~~the architecture constraints memory ~~address which cores are able to access that can be~~directly ~~accessed by the cores~~In order to communicate with cores outside the cluster, ~~we have~~it is necessary to transfer data between clusters through ~~the~~D-NoC with network interfaces.

Comment [Editor12]: Remark: Please check whether the edits retain your intended meaning.

~~In a~~Rx interface exists in the receiving side, ~~there is Rx interface~~to receive data with DMA. ~~We have~~It is necessary to allocate a D-NoC Rx resource and configure it to wait ~~receiving to receive the~~data. ~~One~~A DMA in a network interface contains 256 D-NoC Rx resources. ~~In sending side, we have two~~Two interfaces ~~for users to send data between clusters~~One is, namely a Tx interface and ~~the other is a~~UC interface as explained in Sections 2.1.1 and 2.1.2, ~~are present with respect to the sending side for users to send data between clusters~~The UC is a network processor ~~programmable that is programmed~~to set threads ~~sending to send~~data in DMA. It executes ~~programmed a~~programed pattern and sends data through ~~the~~D-NoC without a PE and an RM. ~~The UC interface provides results in faster data transfer than when compared to that in the~~Tx interface. However, ~~One~~a DMA in a network interface contains only 8 D-NoC UC resources. Both interfaces use ~~a~~DMA engine to access memory and copy data. ~~Regardless~~Irrespective of using whether or not a UC interface ~~or not, we have~~is used, it is necessary to allocate a D-NoC Tx resource and configure it to send data. ~~If we use UC interface, we additionally~~Additionally, it is necessary to allocate and configure ~~a~~D-NoC UC resources ~~if a UC interface is used~~.

4. EVALUATIONS

~~In this~~This section, ~~we have conducted~~involves examining two ~~kind~~types of evaluations. ~~One is, namely a~~D-NoC data transfer evaluation. ~~We explore in which~~latency characteristics of interfaces and memory type. ~~The other is are explored and a~~matrix calculation evaluation. ~~This test shows MPPA-256's that demonstrates the~~parallelization potential and characteristics of the MPPA-256 and its memory access ~~when~~characteristics while dealing with large data.

4.1 D-NoC Data Transfer

In this evaluation, we clarify involves clarifying end-to-end latency by considering the relation among interfaces (Tx or UC), routing on NoC, and memory type (DDR or SMEM). To achieve above purpose, we prepare This is achieved by preparing four routes as shown in Figure 5. These The routes on D-NoC map (Figure 3) contains various connections between routers, namely a direct link, a cross link, and a flying link. In ease of With respect to routes from the IOS routers to the CC routers, transmitted data is allocated in DDR or IO SMEM. The CC has-only includes SMEM as shown in Figure 2. We directly use A low-level library is directly used to transfer data with D-NoC. Transferred The transferred data are correspond to 100 B, 1 KB, 10 KB, 100 KB, and 1 MB. These The buffers are

sequentially allocated in DDR or SRAM (IO SMEM or CC SMEM). Since the The capacity of CC SMEM is corresponds to 2 MB, we assume and thus it is assumed that the appropriate communication buffer size is corresponds to 1 MB. On our assumption Given the assumption, the other memory area is own by corresponds to the application, library, and operating system. In many situations, we have measured end End-to-end latencies are measured 1,000 times and draw in numerous situations as shown in Figures 6, 7, and 8, and boxplots are obtained as depicted in Figures 6, 7, 8, and 9, 10 and 11. Following The following evaluations is are conducted on eMCOS.

Data transfer latencies between IOS and CC is are not influenced by routing. We prepare This involved preparing two interfaces (Tx and UC), three routes (direct link, cross link, and detour route), and two memory location where locations in which the transferred data is allocated. As shown in Figures 6, 7, 8, and 9, end-to-end latencies latency scales

Comment [Editor13]: Remark:
Please check whether the edits retain your
intended meaning.




Figure 5: Four D-NoC routes used for in the evaluation.

linearly exhibit a linear relation with data size and, and there are no significant differences between the three routes do not produce differences in with respect to data transfer latency. This result is important in torus topology NoC because the number of its minimum step is larger than steps exceeds those in the mesh topology's one. We can observe topology. It is observed that queuing in the NoC Routers routers and hardware distance on the NoC is are not dominant factor factors for latency. Transmitting The time taken by transmitting and receiving transaction take much longer time than transactions exceeds those of other transaction. In addition, we Additionally, it is briefly recognize recognized that the speed of UC is much faster than exceeds that of Tx. For The data is arranged as shown in Figures 10 and 11 to facilitate a precise analysis for with respect to the interface and memory location, we arrange data in Figures 10 and 11. In these, In the figures, we accept only cross link the crosslink from IOS to CC5 is accepted because routes do not influence latency. For In order to facilitate intuitive recognition, we arrange two kinds of figures, are arranged, namely a logarithmic axis one and a linear axis one.

In the Tx interface, DDR causes a large increase in latency. The time taken by the DDR takes is twice time as long as that of the IO SMEM as shown in Figure 11. This is caused by due to the memory access speed characteristics of DRAM and SRAM. In the case of the Tx interface, it is necessary for an RM (k1-core) on IOS has to operate the DMA in the IOS network interface of IOS. It This is thought that attributed to the fact that the core is involved in processing is causing this result. Data, The speed of the data transfer latency between CCs is a little faster than exceeds that between IOS and CC. This result indicates that the MPPA-256 is optimized for communication between the CCs.

In With respect to the UC interface, whether the latency is not significantly affected by the location at which the transferred buffer is allocated on (i.e., the DDR or SMEM has not much effect on latency. Same). Similar latency characteristics is are observed in Figures 10 and 11. In the case of the UC interface, an RM (k1-core) on the IOS does not involve a DMA transaction. A micro core in the network interface executes a programmed thread sending data. This evaluation result tells that there is no need to be conscious of suggests that the slow access speed of the DDR is not significant in the case of the UC. As well as In a manner similar to the Tx interface, Data the speed of the data transfer latency between CCs is faster than exceeds that between IOS and CC.

4.2 —Matrix Calculation

In this the evaluation, we clarify matrix calculation time and

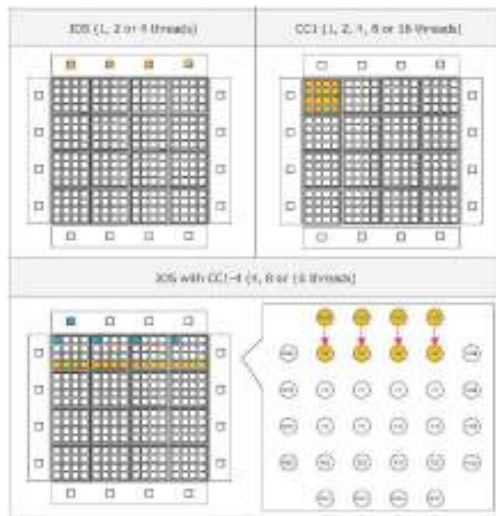


Figure 12: Situations of matrix calculation situations.

parallelization potential of MPPA-256. We have are clarified. Matrix calculations are conducted matrix calculation in IOS and CC. As shown in Figure 12, three Three computing situations is are considered. One is as shown in Figure 12. The first situation involves computing in IOS where four cores are available. To In order to analyze memory access characteristics, we alloentea matrix buffer is allocated in IO DDR and SMEM. Another is The second situation involves computing in CC wherein which 16 cores are available. The other is third situation involves offload-computing by using an IOS and four CCs. Parallelized processing is executed with four CCs cores and there SMEM. Some three SMEMs. A

~~few~~ cores in IOS and CC manage the parallelized transaction. ~~This~~The method ~~is able to~~can handle large data ~~in~~ which one cluster ~~is not sufficient cannot deal~~because buffer capacity is not limited to ~~2MB of~~2 MB in SMEM. Parallelized processing and the total capacity of SMEM ~~is~~are superior to ~~computation in~~IOS or CC. ~~In computations~~. With respect to the IOS, ~~the~~ application can handle large capacity data only in ~~the~~ DDR. However, ~~with~~in this method, ~~we can~~distributed memories are used to deal with large capacity data in SMEM ~~by distributed memories~~. ~~For~~ In order to facilitate faster data transfer, a part of ~~the~~ matrix buffer is ~~parallelly~~transmitted in parallel as shown in Figure 12. ~~To avoid cache coherency trouble~~. Thus, it is necessary for IOS and CC cores ~~must~~to access matrix buffers without cache ~~to avoid cache coherency trouble~~.

~~We analyze matrix~~Matrix calculation time ~~is analyzed~~ with parallelization and memory allocation. ~~In addition, we analyze~~Additionally, the influence of cache ~~is analyzed~~ because cache coherency is an important issue in a many-core system. There are ~~many~~several cases ~~that in which~~ applications must access specific memory space without a cache. ~~Since~~With respect to the given assumptions, maximum buffer size ~~is~~corresponds to 1 MB ~~in our assumption, we prepare, and thus~~ three matrices buffers ~~are prepared~~, and each size ~~is~~corresponds to 314 KB. ~~The matrix~~Matrix A and ~~the~~ matrix B are multiplied, and the result is stored in ~~the~~ matrix C. ~~We set the~~The total of the three matrices ~~to be~~is set as approximately 1 MB.

First, matrix calculation time with ~~the~~ cache in IOS and CC is ~~shown~~depicted in Figure 13. ~~Thanks to cache, there is~~There are almost no ~~difference~~differences between IO DDR, IO SMEM, and CC SMEM ~~due to the cache~~. Furthermore, 128 KB data cache in ~~the~~IOS works well and compensates for the ~~DDR~~ delay ~~of DDR~~. Additionally, ~~we can observe~~it is observed that calculation time scales ~~linearly~~exhibit a linear relation with the number of threads. This ~~is~~corresponds to ideal




Figure 6: Data transfer with Tx from IO DDR to CC.




Figure 7: Data transfer with UC from IO DDR to CC.




Figure 8: Data transfer with Tx from IO SMEM to CC.




Figure 9: Data transfer with UC from IO SMEM to CC.




Figure 10: Data transfer with Tx/UC (logarithmic axis).




Figure 11: Data transfer with Tx/UC (linear axis).

behavior with [respect to](#) parallelization.

Second, matrix calculation time without cache in the IOS and CC is shown in Figure 14. ~~Due to the lack~~The absence of a cache, ~~DDR has increased about~~ results in a fourfold increase in the DDR and a ~~high~~ difference ~~arises~~ with respect to the SMEM ~~has arisen~~. Another notable result is that calculation speed in

CC SMEM ~~is faster than~~exceeds that ~~in~~of the IO SMEM. This characteristic ~~hides~~is hidden in the calculation with the cache. ~~Since~~The computing cores ~~are~~involve the same k1-cores in IOS and CC, and thus it is ~~thought~~considered that the characteristics and physical arrangement of SMEM ~~are affecting~~exert a significant effect. This is ~~an~~an interesting result ~~that there is~~since a ~~high~~ difference that ~~can not~~cannot be ignored. ~~We can~~exists. It is also ~~observe~~observed that calculation time ~~sees linearly~~exhibits a linear relation with the number of threads. ~~The other notable result is that~~Furthermore, the calculation speed without the cache in CC SMEM ~~is faster than~~exceeds that with cache. ~~Although this~~This result is contrary to intuition, and the two factors are conceivable. A small data cache (8 KB) in CC does not ~~function~~adequately ~~work~~, and direct assembly instruction for uncached access optimizes memory access.

Finally, matrix calculation with offload-computing in IOS and CCs is ~~is~~shown in Figure 15. In this case, ~~we suppose that it is assumed with~~respect to the calculation of large matrices ~~calculation that~~ the total capacity ~~is over~~exceeds 1 MB. ~~Because of this assumption, the~~The offloading result is compared with IO DDR (~~cached~~). ~~To~~cached due to the fore-mentioned assumption. An overhead transaction (inverse matB and input matC) ~~is performed to~~ offload a part of calculation ~~on~~on the four CCs and aggregate calculation result, ~~overhead transaction (inverse matB and input matC) occurred. They produce~~. This produces a constant overhead ~~regardless~~irrespective of the number of threads as shown in Figure 15. However, the speed involved in offloading result ~~is faster than~~exceeds that of IO DDR (~~cached~~). This ~~cached~~. The result ~~proves~~indicates several important facts. ~~One is that~~First, D-NoC data transfer ~~produce~~produces little overhead latency. ~~The other is that~~Second, the speed of DMA memory access to DDR ~~is much faster than an~~exceeds that of RM (k1-core)'s memory access. In the offloading case, a DMA accesses matrix buffers on DDR and ~~transfer them~~transfers the buffers from IO

DDR to each CC SMEM. ~~Then~~Subsequently, PEs in the CC access matrix

~~buffers for~~buffer the calculation without cache. ~~Since the~~The overhead of data transfer and DMA memory access is small, and thus parallel data transmission and distributed memory are practical in the case of MPPA-256. The impact of offloading ~~becomes larger~~increases when the matrix size is large as shown in Figure 16. ~~Since only~~Only a part of matrix is allocated in CC, ~~we can~~and thus it is possible to handle larger matrices buffers. ~~We prepare~~Additionally, 640 KB matrices are prepared, and ~~evaluate~~matrix ~~calculation~~calculations are evaluated with offload-computing. ~~Compared to result of 314 KB matrices in Figure 15~~The speed of the offloading result ~~is much faster than~~exceeds that of IO DDR result ~~with respect to the 314 KB matrices~~ in Figure 15.

5. RELATED WORK

In this

This section, ~~we compare~~compares many-core platforms ~~to additional platforms~~and ~~discuss~~discusses previous ~~work of~~studies related to multi/many cores.

Comment [Editor14]: Remark: Note that this section forms the latter part of the introduction section. Hence, please consider moving this to the last part of the introduction section.

~~In recent years,~~ Recently, studies indicate that the single core processors are characterized by limited computation performance ~~of~~. Pollack stated ~~that~~ a single core ~~processor~~ is constantly coming to its limit. Pollack has said inefficiency of a single core is ~~inefficient~~ [23] and ~~that~~ Moore's law [19] ~~has become unstable. To is no longer applicable.~~ Therefore, extant CPUs are not sufficient to satisfy increasing a demand of computation. ~~CPU is not sufficient,~~ demands. Many other platforms including many-core are developed and researched ~~nowadays~~ by current studies.

Table 1 summarizes ~~the~~ features of many-core platforms with that of other platforms. For instance, ~~the~~ GPU is a powerful device to enhance computing performance. ~~In, and it has great potential in specific area (areas (for e.g., image processing, and learning), it has great potential).~~ However, it is mainly used for a specific purpose, and its reliability is not suitable for real-time systems. It is difficult to use a GPU for a global purpose and ~~to~~ guarantee ~~efits~~ reliability due to ~~the~~ GPU architecture. ~~For~~ Many-core is significantly superior to GPU with respect to a global purpose and multiple instructions, ~~many-core is much superior to GPU. In addition,~~ Additionally, it is commonly known that many-core ~~has~~ involves a reasonable power consumption. ~~In contrast, the~~ GPU consumes a ~~lot~~ significant amount of power and generates ~~much~~ considerable heat. This is a critical problem for embedded systems. ~~DSP~~ Furthermore, DSPs and ~~FPGA~~ FPGAs are also high-performance devices ~~when~~ compared to ~~CPU~~ CPUs. They are efficient in ~~a point~~ terms of power consumption. A DSP is often used for real-time ~~system~~ systems, and ~~FPGA~~ FPGAs ~~guarantee~~ FPGAs guarantee reliability and efficient processing. They




Figure 15: Matrix ~~calculation~~calculations with offload-computing (314 KB matrix x 3).



Figure 16: Matrix ~~calculation~~calculations with offload-computing (640 KB matrix x 3).

are suitable for time-critical computing [10]. However, ~~DSP~~DSPs cannot be used for global purpose and ~~programming-FPGA-is~~programming. Furthermore, FPGAs are difficult for software developers. ~~Their~~since their software model is ~~much~~significantly different from that of CPU and is not a substitute for CPU. Many-core platforms ~~have a potential to take the place of~~can potentially replace single/multi-core CPU. ~~Ease as they possess ease of~~programmingprogramming and scalability with high acceleration-is highlighted.

Based on the ~~above~~fore-mentioned background, ~~several extant studies examined~~ real-time applications on many-core platforms ~~has received significant attention in recent years.~~ Many ~~commercial-off-the-shelf (COTS)-multi-core~~ multicore components are developed and released by several vendors. (e.g., Kalray's the Multi-Purpose ~~Pro-cessing~~Processing Array

(MPPA) 256 [12], [11], [15], Intel's Single-chip Cloud Computer (SCC) [1],[3], Tilera's Tile64 [2], and Intel's Xeon Phi

[7] [8]). The ~~present study focuses on the~~ Kalray MPPA-256 ~~that~~ is designed for real-time embedded applications ~~and the target of this paper.~~ Kalray

5. A [12], [11], [15] ~~has~~presented clustered many-core architectures on ~~the~~NoC. ~~(Precise~~The precise hardware model is described in Section 2.1. ~~It~~. This is ~~often~~typically accepted ~~for a target of~~with respect to many-core platforms, and ~~various~~the model is used in several previous ~~work has considered this~~

~~model~~studies [22], [4], [6], [21], [TBD]

6. ~~CONCLUSION~~CONCLUSIONS

[TBD]

Table 1: ~~Comparison~~A comparison of Many-core to CPU, GPU, DSP, and FPGA

	performance	power/heat	reliability	real-time	software development	costs	multiple instruction
CPU		A	✓	✓	✓	✓	A
GPU	✓		A		A	✓	
DSP	A	A	✓	✓	A	✓	
FPGA	✓	A	✓	A		A	
Many-core	✓	✓	✓	✓	✓	A	✓